

Homework 2

CS 5700 - Database Systems

17.28. A file has $r = 20,000$ STUDENT records of fixed length. Each record has the following fields: Name (30 bytes), Ssn (9 bytes), Address (40 bytes), PHONE (10 bytes), Birth_date (8 bytes), Sex (1 byte), Major_dept_code (4 bytes), Minor_dept_code (4 bytes), Class_code (4 bytes, integer), and Degree_program (3 bytes). An additional byte is used as a deletion marker. The file is stored on the disk whose parameters are given in Exercise 17.27.

- a. Calculate the record size R in bytes.

$$R = 30 + 9 + 40 + 10 + 8 + 1 + 4 + 4 + 4 + 3 + 1$$

$$R = 114 \text{ bytes}$$

- b. Calculate the blocking factor bfr and the number of file blocks b , assuming an unspanned organization.

$$bfr = \text{floor}(\text{blockSize} / R)$$

$$bfr = \text{floor}(512 / 114)$$

$$bfr = 4$$

- c. Calculate the average time it takes to find a record by doing a linear search on the file if (i) the file blocks are stored contiguously, and double buffering is used; (ii) the file blocks are not stored contiguously.

i. $\text{initial seek} + (b / 2) - 1 * \text{sequential seek}$, where b is the number of blocks

$$= 30 + ((20,000 / 4) / 2) - 1 * 1$$
$$= 2,529 \text{ ms}$$

ii. $\text{random seek} * (b / 2)$, where b is the number of blocks

$$= 30 * ((20,000 / 4) / 2)$$
$$= 75,000 \text{ ms}$$

d. Assume that the file is ordered by Ssn; by doing a binary search, calculate the time it takes to search for a record given its Ssn value.

$$\begin{aligned} & \text{random seek} * \log(b), \text{ where } b \text{ is the number of blocks} \\ & = 30 * \log_2(2,500) \\ & = \text{approx. } 338.631 \text{ ms} \end{aligned}$$

18.18. Consider a disk with block size $B = 512$ bytes. A block pointer is $P = 6$ bytes long, and a record pointer is $P_R = 7$ bytes long. A file has $r = 30,000$ EMPLOYEE records of fixed length. Each record has the following fields: Name (30 bytes), Ssn (9 bytes), Department_code (9 bytes), Address (40 bytes), Phone (10 bytes), Birth_date (8 bytes), Sex (1 byte), Job_code (4 bytes), and Salary (4 bytes, real number). An additional byte is used as a deletion marker.

a. Calculate the record size R in bytes.

$$\begin{aligned} R &= 30 + 9 + 9 + 40 + 10 + 8 + 1 + 4 + 1 \\ R &= 112 \text{ bytes} \end{aligned}$$

b. Calculate the blocking factor bfr and the number of file blocks b , assuming an unspanned organization.

$$\begin{aligned} bfr &= \text{floor}(\text{blockSize} / R) \\ bfr &= \text{floor}(512 / 112) \\ bfr &= 4 \end{aligned}$$

c. Suppose that the file is ordered by the key field Ssn and we want to construct a primary index on Ssn. Calculate (i) the index blocking factor bfr_i (which is also the index fan-out fo); (ii) the number of first-level index entries and the number of first-level index blocks; (iii) the number of levels needed if we make it into a multilevel index; (iv) the total number of blocks required by the multilevel index; and (v) the number of block accesses needed to search for and retrieve a record from the file—given its Ssn value—using the primary index.

$$\begin{aligned} \text{i. } R_i &= \text{size of ssn} + \text{size of block pointer} \\ &= 9 + 6 \\ &= 15 \text{ bytes} \end{aligned}$$

- $$\begin{aligned} \text{bfri} &= \text{floor}(\text{blockSize} / \text{Ri}) \\ &= \text{floor}(512 / 15) \\ &= 34 = \text{fo} \end{aligned}$$
- ii. $\text{first-level entries} = \text{number of entries} / \text{bfr}$
 $= 30,000 / 4$
 $= 7,500 \text{ entries}$
 $\text{first-level blocks} = \text{ceil}(\text{first-level entries} / \text{bfri})$
 $= \text{ceil}(7,500 / 34)$
 $= 221 \text{ blocks}$
- iii. $\text{number of levels} = \text{ceil}(\log_{\text{fo}}(\text{first-level entries}))$
 $= \text{ceil}(\log_{34}(7500))$
 $= 3 \text{ levels}$
- iv. $\text{total index blocks} = \text{level three blocks} + \text{level two blocks} + \text{level one blocks}$
 $= 1 + \text{ceil}(193 / 34) + 193$
 $= 1 + 6 + 193$
 $= 200 \text{ blocks}$
- v. $\text{total blocks retrieved for primary lookup} = \log_2(\text{number of first-level blocks}) + 1$
 $= 8 + 1$
 $= 9 \text{ blocks retrieved}$

d. Suppose that the file is not ordered by the key field Ssn and we want to construct a secondary index on Ssn. Repeat the previous exercise (part c) for the secondary index and compare with the primary index.

- i. $\text{Ri} = \text{size of ssn} + \text{size of record pointer}$
 $= 9 + 7$
 $= 16 \text{ bytes}$
 $\text{bfri} = \text{floor}(\text{blockSize} / \text{Ri})$
 $= \text{floor}(512 / 16)$
 $= 32 = \text{fo}$
- ii. $\text{first-level entries} = \text{number of entries}$
 $= 30,000 \text{ entries}$

- first-level blocks = $\text{ceil}(\text{first-level entries} / \text{bfri})$
= $\text{ceil}(30,000 / 32)$
= 938 blocks
- iii. number of levels = $\text{ceil}(\log_{\text{fo}}(\text{first-level entries}))$
= $\text{ceil}(\log_{32}(30,000))$
= 3 levels
- iv. total index blocks = level three blocks + level two blocks + level one blocks
= $1 + \text{ceil}(938 / 32) + 938$
= $1 + 30 + 770$
= 801 blocks
- v. total blocks retrieved for primary lookup = $\log_2(\text{number of first-level blocks}) + 1$
= $10 + 1$
= 11 blocks retrieved

e. Suppose that the file is not ordered by the nonkey field `Department_code` and we want to construct a secondary index on `Department_code`, using option 3 of Section 18.1.3, with an extra level of indirection that stores record pointers. Assume there are 1,000 distinct values of `Department_code` and that the `EMPLOYEE` records are evenly distributed among these values. Calculate (i) the index blocking factor `bfri` (which is also the index fan-out `fo`); (ii) the number of blocks needed by the level of indirection that stores record pointers; (iii) the number of first-level index entries and the number of first-level index blocks; (iv) the number of levels needed if we make it into a multilevel index; (v) the total number of blocks required by the multilevel index and the blocks used in the extra level of indirection; and (vi) the approximate number of block accesses needed to search for and retrieve all records in the file that have a specific `Department_code` value, using the index.

- i. $R_{\text{indirection}} = \text{size of record pointer}$
= 7 bytes
 $\text{bfri}_{\text{indirection}} = \text{floor}(\text{blockSize} / R_i)$
= $\text{floor}(512 / 7)$
= 73 = `fo`

- ii. records per unique value = $30,000 / 1,000$
= 30 distinct records
blocks per value = 1, since 30 records > 73 bfrindirection
blocks for indirection layer = number of distinct values * blocks per value
= $1000 * 1$
= 1000 blocks
- iii. R_i = size of Department_code + size of block pointer
= $9 + 6$
= 15 bytes
 $bfri = \text{floor}(\text{blockSize} / R_i)$
= $\text{floor}(512 / 15)$
= $34 = fo$
number of first-level entries = number of distinct values
= 1000 entries
number of first-level blocks = $\text{ceil}(\text{number of first-level entries} / bfri)$
= $\text{ceil}(1000 / 34)$
= 30 blocks
- iv. number of levels = $\text{ceil}(\log_{fo}(\text{first-level entries}))$
= $\text{ceil}(\log_{34}(1,000))$
= 2 levels
- v. total multilevel blocks = second-level blocks + first-level blocks + indirection layer
= $1 + 30 + 1000$
= 1031 blocks
- vi. approximate total block accesses
= $\log_2(\text{first-level blocks})$
+ blocks per indirection entry
+ avg entries per indirection entry
= $5 + 1 + 30$
= 36 blocks

f. Suppose that the file is ordered by the nonkey field `Department_code` and we want to construct a clustering index on `Department_code` that uses block anchors (every new value of `Department_code` starts at the beginning of a new block). Assume there are 1,000 distinct values of `Department_code` and that the `EMPLOYEE` records are evenly distributed among these values. Calculate (i) the index blocking factor bfr_i (which is also the index fan-out fo); (ii) the number of first-level index entries and the number of first-level index blocks; (iii) the number of levels needed if we make it into a multilevel index; (iv) the total number of blocks required by the multilevel index; and (v) the number of block accesses needed to search for and retrieve all records in the file that have a specific `Department_code` value, using the clustering index (assume that multiple blocks in a cluster are contiguous).

- i. $R_i = \text{size of } Department_code + \text{size of block pointer}$
 $= 9 + 6$
 $= 15 \text{ bytes}$
 $bfr_i = \text{floor}(\text{blockSize} / R_i)$
 $= \text{floor}(512 / 15)$
 $= 34 = fo$
- ii. $\text{first-level entries} = \text{number of distinct entries}$
 $= 1,000 \text{ entries}$
 $\text{first-level blocks} = \text{ceil}(\text{first-level entries} / bfr_i)$
 $= \text{ceil}(1,000 / 34)$
 $= 30 \text{ blocks}$
- iii. $\text{number of levels} = \text{ceil}(\log_{fo}(\text{first-level entries}))$
 $= \text{ceil}(\log_{34}(1000))$
 $= 2 \text{ levels}$
- iv. $\text{total index blocks} = \text{level two blocks} + \text{level one blocks}$
 $= 1 + 30$
 $= 31 \text{ blocks}$
- v. $\text{size of clustering group} = \text{total records} / \text{number of distinct entries}$
 $= 30,000 / 1,000$
 $= 30 \text{ blocks}$

total blocks retrieved for primary lookup
= $\log_2(\text{number of first-level blocks}) + \text{size of clustering group}$
= $5 + 30$
= 39 blocks retrieved

g. Suppose that the file is not ordered by the key field Ssn and we want to construct a B⁺-tree access structure (index) on Ssn. Calculate (i) the orders p and p_{leaf} of the B⁺-tree; (ii) the number of leaf-level blocks needed if blocks are approximately 69 percent full (rounded up for convenience); (iii) the number of levels needed if internal nodes are also 69 percent full (rounded up for convenience); (iv) the total number of blocks required by the B⁺-tree; and (v) the number of block accesses needed to search for and retrieve a record from the file—given its Ssn value—using the B⁺-tree.

- i. $(p * \text{block pointer size}) + ((p - 1) * \text{size of ssn}) \leq \text{block size}$
 $6p + 9p - 9 \leq 512$
 $15p \leq 521$
 $p \leq 34.73 \text{ bytes}$
 $p = 34$
 $(p_{\text{leaf}} * (\text{size of ssn} + \text{size of record pointer})) + \text{size of next pointer} \leq \text{block size}$
 $(9 + 7)p_{\text{leaf}} + 6 \leq 512$
 $16p_{\text{leaf}} \leq 506$
 $p_{\text{leaf}} \leq 31.625 \text{ bytes}$
 $p_{\text{leaf}} = 31$
- ii. number of leaf-level blocks at 69% capacity = number of records / (0.69 * p_{leaf})
= $30,000 / \text{floor}(0.69 * 31)$
= $\text{ceil}(30,000 / 21)$
= 1,429 leaf nodes
- iii. fop = round(0.69 * p)
= 23 pointers
number of levels = $\text{ceil}(\log_{\text{fop}}(\text{number of leaf nodes}))$
= $\text{ceil}(\log_{23}(1429))$
= 3 levels

- iv. total blocks = level-three blocks + level-two blocks + level-one blocks + leaf blocks
= $1 + 23 + 23^2 + 1429$
= 1982 blocks
- v. number of block accesses
= level-three lookup
+ level-two lookup
+ level-one lookup
+ leaf lookup
+ record lookup
= $1 + 1 + 1 + 1 + 1$
= 5 blocks

- 19.13. Consider queries Q1, Q8, Q1B, and Q4 in Chapter 4 and Q27 in Chapter 5.
a. Draw at least two query trees that can represent each of these queries. Under what circumstances would you use each of your query trees?

SEE ATTACHMENT FOR TREES

Q1:

```
SELECT Fname, Lname, Address
FROM EMPLOYEE, DEPARTMENT
WHERE Dname='Research'AND Dnumber=Dno;
```

Q8:

```
SELECT E.Fname, E.Lname, S.Fname, S.Lname
FROM EMPLOYEE AS E, EMPLOYEE AS S
WHERE E.Super_ssn=S.Ssn;
```


Q1B:

```
SELECT E.Fname, E.LName, E.Address  
FROM EMPLOYEE E, DEPARTMENT D  
WHERE D.DName='Research'AND D.Dnumber=E.Dno;
```

Q4:

```
(SELECT DISTINCT Pnumber  
FROM PROJECT, DEPARTMENT, EMPLOYEE  
WHERE Dnum=Dnumber AND Mgr_ssn=Ssn  
AND Lname='Smith' )  
UNION  
( SELECT DISTINCT Pnumber  
FROM PROJECT, WORKS_ON, EMPLOYEE  
WHERE Pnumber=Pno AND Essn=Ssn  
AND Lname='Smith' );
```

Q27:

```
SELECT Pnumber, Pname, COUNT (*)  
FROM PROJECT, WORKS_ON, EMPLOYEE  
WHERE Pnumber=Pno AND Ssn=Essn AND Dno=5  
GROUP BY Pnumber, Pname;
```

b. Draw the initial query tree for each of these queries, and then show how the query tree is optimized by the algorithm outlined in Section 19.7.

SEE ATTACHMENT

c. For each query, compare your own query trees of part (a) and the initial and final query trees of part (b).

Q1: My second query tree is similar to the final tree of part b, only I did not move the projections as far down the tree as possible.

Q8: My second query tree is the same as the tree produced by the algorithm.

Q1B: My first query tree is the same as the tree produced by the algorithm, my second tree is arguably better, but it relies on an assumption that may or may not be true.

Q4: My second query tree is similar to the final query tree for Q4, but I left off a few projections here and there.

19.17 Can a nondense index be used in the implementation of an aggregate operator? Why or why not?

It can, but extra care must be taken to insure that the proper data is retrieved. For MAX and MIN aggregates, we must be sure that the smallest and largest values are stored in our index, such as how a B+-tree would store them. For AVG or SUM aggregates we must maintain a count of the number of values per index entry so that we may use the actual total number of elements in these functions. This is also similar to how a COUNT aggregate would function on a sparse index.

19.22 Compare the cost of two different query plans for the following query:

$\sigma_{\text{Salary} > 40000}(\text{EMPLOYEE} \bowtie_{\text{Dno}=\text{Dnumber}} \text{DEPARTMENT})$

Use the database statistics in Figure 19.8.

SEE ATTACHMENT FOR TREES